

TABLE OF CONTENTS

CHAPTER 1: Data Streaming and Analysis	1
1.0 Introduction.....	2
1.1 Definition of Data Streaming and Analysis	2
1.1.1 What Is Data Streaming?.....	2
1.1.2 What Is Data Analysis in a Streaming Context?.....	2
1.1.3 Batch Processing versus Stream Processing	3
1.2 Key Concepts of Data Streaming and Analysis	4
1.2.1 Bounded and Unbounded Data.....	4
1.2.2 How Modern Systems Handle Both.....	5
1.2.3 Core Characteristics of Stream Processing.....	5
1.2.4 The Stream-Processing Workflow.....	6
1.2.5 An End-to-End View	9
1.3 Use Cases and Applications in Various Domains.....	10
1.3.1 Domain-wise Use Cases.....	10
1.4 Real-World Examples	11
1.4.1 Netflix and Amazon.....	11
1.4.2 Twitter (X) and Tesla	12
1.4.3 Hospitals — Real-Time Patient Monitoring.....	12
1.5 Potential Benefits of Data Streaming and Analysis	13
1.6 Tools and Technologies	13
1.6.1 Key Tools at a Glance.....	14
1.7 Comparison of Tools and Technologies	15
1.8 Strengths and Limitations of Tools and Technologies.....	17
1.9 Module Summary	17
1.10 Key Terms	18
1.11 Review Questions	19
Short Answer	19
Essay / Discussion.....	19
1.12 References	19

CHAPTER 2: Streaming Data Processing	20
2.1 Introduction to Streaming Data Processing	20
2.1.1 Batch Processing vs. Streaming Processing.....	20
2.1.2 Principles of Streaming Data Processing	21
2.1.3 Key Challenges in Streaming Data Processing.....	22
2.2 Streaming System Architecture and Distributed Foundations	23
2.2.1 Role of Distributed Systems in Streaming Data Processing	23
2.2.2 Stream Ingestion and Event Processing — Core Concepts	24
2.2.3 Generic Architecture of a Streaming Pipeline	24
2.3 Apache Kafka: Architecture and Working.....	25
2.3.1 Introduction to Apache Kafka.....	26
2.3.2 Kafka Architecture.....	26
2.3.3 Topics and Partitions	27
2.3.4 Producers	28
2.3.5 Consumers and Consumer Groups	28
2.3.6 Brokers and the Kafka Cluster	29
2.3.7 Offsets and Message Ordering	30
2.3.8 Kafka Data Flow.....	30
2.3.9 Reliability and Fault Tolerance in Kafka	31
2.3.10 Kafka Working Mechanism — Step-by-Step Operational Flow	31
2.4 Apache Storm: Architecture and Working.....	32
2.4.1 Introduction to Apache Storm	32
2.4.2 Storm Architecture.....	34
2.4.3 Topologies, Spouts, and Bolts	36
2.4.4 Stream Grouping Strategies	38
2.4.5 Storm Data Flow.....	40
2.4.6 Reliability and Fault Tolerance in Storm	40
2.4.7 Storm Working Mechanism — Step-by-Step Operational Flow.....	42
2.4.8 Common Design Considerations for Storm Topologies	43
2.4.9 Storm Processing Guarantees at a Glance.....	44
2.4.10 Section 2.4 — Key Terminology Recap.....	45
2.5 Comparative Analysis: Apache Kafka vs. Apache Storm	46

2.5.1 Core Comparison.....	46
2.5.2 Programming Model Comparison	47
2.5.3 Common Misconceptions	48
2.5.4 Complementary Roles in a Streaming Architecture.....	48
2.5.5 When to Use Kafka Alone, Storm Alone, or Both Together	49
2.5.6 Fault-Tolerance Mechanisms Compared	50
2.5.7 Operational Considerations	51
2.5.8 A Practical Decision Checklist.....	52
2.5.9 Summary of the Comparison	52
2.6 Handling and Managing Real-Time Data Streams	53
2.6.1 Real-Time vs. Micro-Batch Processing.....	53
2.6.2 Event Time vs. Processing Time Semantics.....	54
2.6.3 Introduction to Windowing in Stream Processing.....	55
2.6.4 Introduction to State Management in Streaming Systems	56
2.6.5 Ensuring Data Ordering and Consistency in Streams.....	57
2.6.6 Section 2.6 — Key Terminology Recap	59
2.7 Building Scalable and Fault-Tolerant Streaming Pipelines.....	60
2.7.1 Principles of Scalability in Streaming Systems.....	60
2.7.2 Partitioning and Load-Balancing Techniques	61
2.7.3 Replication and Fault-Tolerance Techniques.....	62
2.7.4 Checkpointing and Recovery Mechanisms	62
2.7.5 Illustrative Scenario: Designing an End-to-End Scalable Pipeline.....	63
2.7.6 Section 2.7 — Key Terminology Recap	65
2.8 Real-World Domains for Streaming Applications.....	65
2.8.1 IoT and Sensor Data Streaming	65
2.8.2 Banking and Financial Fraud Detection.....	65
2.8.3 Social Media Analytics	66
2.8.4 Transportation and Logistics.....	66
2.8.5 Healthcare Monitoring.....	66
2.8.6 Cybersecurity (Intrusion and Anomaly Detection).....	66
2.9 Case Studies on Real-World Data Streaming Applications	66
2.9.1 Case Study: Kafka-Based Log Aggregation in an E-Commerce Platform....	67

2.9.2 Case Study: Real-Time IoT Sensor Pipeline Using Kafka and Storm	67
2.9.3 Case Study: Smart City Traffic Monitoring Using Streaming Data	67
2.9.4 Summary Table of Domains and Case Studies.....	67
2.9.5 Common Patterns Across Domains and Case Studies.....	68
2.10 Unit Summary	69
2.11 Key Terms.....	69
2.12 Review Questions	70
2.13 Multiple Choice Questions	70
2.14 Application/Problem-Solving Questions	71
2.15 Practical/Laboratory Activities	72
2.16 References	72
CHAPTER 3: Data Transformation and Analysis	73
3.1 Introduction to Data Transformation and Analysis.....	73
3.1.1 Concept and Need for Data Transformation in Streaming Systems	73
3.1.2 Principles of Data Transformation and Analysis.....	73
3.1.3 From Streaming Ingestion to Real-Time Analytics	74
3.2 Apache Spark: Architecture and Working.....	75
3.2.1 Introduction to Apache Spark.....	75
3.2.2 Spark Architecture	77
3.2.3 Resilient Distributed Datasets (RDDs) and DataFrames	78
3.2.4 Spark Streaming and Structured Streaming	79
3.2.5 Spark Data Flow	80
3.2.6 Fault Tolerance in Spark	81
3.2.7 Spark Working Mechanism — Step-by-Step Operational Flow	82
3.3 Apache Flink: Architecture and Working.....	84
3.3.1 Introduction to Apache Flink.....	84
3.3.2 Flink Architecture.....	86
3.3.3 Flink Streaming Dataflow Model	87
3.3.4 Windowing in Flink.....	88
3.3.5 State Management in Flink.....	90
3.3.6 Fault Tolerance — Checkpointing and Savepoints in Flink.....	91
3.3.7 Flink Working Mechanism — Step-by-Step Operational Flow	91

Flink Performance and Scalability Considerations	93
3.4 Comparative Analysis: Apache Spark vs. Apache Flink.....	93
Common Pitfalls When Choosing Between Spark and Flink	96
A Practical Decision Checklist.....	96
3.5 Data Cleaning Techniques for Streaming Data.....	97
3.5.1 Handling Missing and Incomplete Data	97
3.5.2 Handling Duplicate and Noisy Data.....	98
3.5.3 Data Validation and Consistency Checks	99
3.6 Data Aggregation and Real-Time Analysis Techniques.....	100
3.6.1 Aggregation Concepts over Streaming Windows.....	100
3.6.2 Real-Time Analytics Techniques	101
3.6.3 Illustrative Scenario: Aggregating Streaming Sensor Data	102
3.7 Real-Time Dashboard Development and Visualization.....	103
3.7.1 Principles of Real-Time Dashboard Design	103
3.7.2 Techniques for Developing Real-Time Dashboards.....	105
3.7.3 Visualization Techniques for Streaming Data	105
3.7.4 Overview of Real-Time Visualization Tools.....	106
3.8 Challenges and Solutions in Data Transformation and Analysis.....	107
3.8.1 Latency vs. Accuracy Trade-Offs.....	107
3.8.2 Handling Late and Out-of-Order Data	107
3.8.3 Scalability Challenges in Analytics Pipelines.....	108
3.8.4 Data Quality and Governance Challenges	109
3.9 Real-World Applications and Case Studies.....	110
3.9.1 Case Study: Real-Time Fraud Detection Dashboard in Banking (Spark) .	110
3.9.2 Case Study: Flink-Based Real-Time Recommendation Analytics in E-Commerce.....	111
3.9.3 Case Study: Industrial Monitoring Dashboard for Predictive Maintenance	111
3.9.4 Case Study: Real-Time Patient Monitoring Analytics in Healthcare.....	112
Summary Table of Unit 3 Case Studies.....	112
Common Patterns Across the Unit 3 Case Studies	112
3.10 Unit Summary	113
3.11 Key Terms.....	114

3.12 Review Questions	115
3.13 Multiple Choice Questions	116
3.14 Application/Problem-Solving Questions	117
3.15 Practical/Laboratory Activities	118
3.16 References	118
CHAPTER 4: Real-Time Data Processing	119
4.1 Introduction to Real-Time Data Processing	119
4.1.1 What is Real-Time Data Processing?	119
4.1.2 Key Characteristics.....	119
4.1.3 A Typical Real-Time Pipeline	120
4.1.4 Why Real-Time Processing Matters	120
4.2 Real-Time Data Processing vs Batch Processing.....	121
4.2.1 Batch Processing.....	121
4.2.2 Stream (Real-Time) Processing.....	121
4.2.3 Comparative Summary	121
4.2.4 Choosing Between Batch and Streaming	121
4.3 Real-Time Data Processing Tools and Frameworks	122
4.3.1 Message Brokers and Ingestion Systems	122
4.3.2 Stream-Processing Engines.....	122
4.3.3 Serving, Storage and Visualisation	122
4.4 Real-Time Data Processing with Apache Kafka	123
4.4.1 Core Concepts.....	123
4.4.2 Kafka Architecture.....	123
4.4.3 Setting Up a Topic.....	123
4.4.4 A Simple Kafka Producer (Java).....	124
4.4.5 A Simple Kafka Consumer (Java)	124
4.4.6 Delivery Guarantees	124
4.5 Processing Real-Time Data Streams with Apache Flink	125
4.5.1 Why Flink?	125
4.5.2 Core Architectural Concepts	125
4.5.3 Event Time, Processing Time, and Windows.....	126
4.5.4 A DataStream API Example (Java)	126

4.5.5 Fault Tolerance: Checkpoints and Savepoints.....	127
4.6 Implementing Real-Time Data Processing Use Cases	127
4.6.1 A General Design Process.....	127
4.6.2 Case Study 1 – Real-Time Fraud Detection in Banking.....	128
4.6.3 Case Study 2 – Real-Time Recommendation Engine	128
4.6.4 Common Implementation Pitfalls.....	128
4.7 Real-Time Analytics for Streaming Data.....	129
4.7.1 Common Analytical Patterns	129
4.7.2 Architecture for Real-Time Dashboards	129
4.8 Real-Time Decision Making with Streaming Data	129
4.8.1 Complex Event Processing (CEP)	129
4.8.2 Rule Engines and Streaming ML.....	130
4.8.3 Human-in-the-Loop vs Fully Automated Decisions	130
4.9 Real-Time Monitoring and Alerting.....	130
4.9.1 What to Monitor	130
4.9.2 Monitoring Architecture	130
4.9.3 Designing Effective Alerts	131
4.10 Unit Summary	131
4.11 Key Terms.....	132
4.12 Review Questions	132
4.13 References	133
CHAPTER 5: Data Visualization and Analytics.....	134
5.1 Introduction to Data Visualization and Analytics	134
5.1.1 What Is Data Visualization?	134
5.1.2 Why Data Visualization and Analytics Matter	134
5.1.3 The Data-to-Decision Pipeline	135
5.1.4 Descriptive, Diagnostic, Predictive, and Prescriptive Analytics	136
5.1.5 Types of Data and Their Visual Implications.....	136
5.1.6 Analytics Maturity and the Rise of Real-Time Visualization	137
5.2 Types of Data Visualization Tools	137
5.2.1 The Business Intelligence and Visualization Tool Landscape.....	137
5.2.2 Tableau.....	138

5.2.3 Microsoft Power BI.....	138
5.2.4 Other Prominent Platforms.....	139
5.2.5 Open-Source and Programmatic Visualization	139
5.2.6 Spreadsheet Tools	140
5.2.7 Choosing a Tool: A Framework.....	140
5.2.8 Chapter Summary	141
5.3 Understanding the Visualization and Analysis of Data Streaming	141
5.3.1 Defining Streaming Data	141
5.3.2 Batch versus Stream Processing.....	141
5.3.3 Core Concepts in Stream Processing	142
5.3.4 Streaming Platforms and Architecture Patterns	143
5.3.5 Two Dominant Processing Paradigms: Lambda and Kappa.....	144
5.3.6 Characteristics That Distinguish Streaming Visualization	144
5.4 Data Visualization for Real-Time Data Streaming	144
5.4.1 What Makes Real-Time Visualization Different	144
5.4.2 Latency Budgets and the Meaning of "Real Time"	145
5.4.3 Chart Types Suited to Streaming Data	146
5.4.4 Design Techniques for Continuously Updating Displays.....	146
5.4.5 Interactivity in Real-Time Dashboards	147
5.4.6 Common Pitfalls in Real-Time Visualization	148
5.4.7 Chapter Summary	148
5.5 Best Practices for Data Visualization and Analysis	148
5.5.1 Why Best Practices Matter	148
5.5.2 Know Your Audience.....	148
5.5.3 Choose the Right Chart for the Data and the Message	149
5.5.4 Minimize Chartjunk and Maximize the Data-Ink Ratio.....	149
5.5.5 Use Color with Purpose	150
5.5.6 Provide Context.....	150
5.5.7 Ensure Accuracy and Avoid Distortion	150
5.5.8 Design for the Delivery Context.....	151
5.5.9 Validate and Iterate.....	151
5.6 Implementing Real-Time Data Visualization and Analytics Use Cases	151

5.6.1 From Principles to Practice.....	151
5.6.2 Case Study: Real-Time Equipment Monitoring and Predictive Maintenance	152
5.6.3 Case Study: Financial Market Surveillance.....	153
5.6.4 Case Study: Application and Infrastructure Monitoring (Observability).....	153
5.6.5 Case Study: Logistics and Fleet Tracking.....	154
5.6.6 Cross-Cutting Implementation Lessons	154
5.7 Exploratory Data Analysis for Streaming Data	155
5.7.1 What Exploratory Data Analysis Means in a Streaming Context	155
5.7.2 Windowed Exploration	155
5.7.3 Distribution and Shape Checks.....	156
5.7.4 Outlier and Anomaly Detection in EDA.....	157
5.7.5 Correlation and Relationship Discovery Across Streams.....	157
5.7.6 Seasonality and Trend Decomposition	157
5.7.7 Visual Techniques for Streaming EDA	158
5.8 Statistical Analysis for Streaming Data	158
5.8.1 The Role of Statistics in Streaming Analytics	158
5.8.2 Statistical Process Control	159
5.8.3 Hypothesis Testing on Streaming Data.....	160
5.8.4 Change Point and Drift Detection.....	160
5.8.5 Time Series Models for Streaming Forecasting.....	161
5.8.6 Non-Stationarity: The Central Statistical Challenge of Streaming Data ...	161
5.9 Machine Learning for Streaming Data.....	162
5.9.1 Why Streaming Data Requires a Distinct Machine Learning Approach ...	162
5.9.2 Online (Incremental) Learning Algorithms	162
5.9.3 Streaming Anomaly Detection with Machine Learning.....	163
5.9.4 Concept Drift: Definition and Detection.....	163
5.9.5 The Online Machine Learning Pipeline	164
5.9.6 Visualizing Machine Learning Output in a Streaming Context	165
5.9.7 Practical Considerations and Limitations.....	165
5.10 Data Storytelling with Real-Time Data Visualizations	166
5.10.1 What Is Data Storytelling?	166

5.10.2 Why Storytelling Matters More, Not Less, for Real-Time Data	166
5.10.3 The Narrative Arc Applied to Streaming Data.....	166
5.10.4 Techniques for Real-Time Data Storytelling	167
5.10.5 Balancing Narrative and Rigor	168
5.10.6 Storytelling for Different Real-Time Audiences	168
5.10.7 Bringing It Together: Storytelling as the Capstone of Visual Analytics..	169
5.11 Unit Summary	169
5.12 Key Terms	170
5.13 Review Questions	170
Appendix: Laboratory Experiments, Assignments, Case Studies and Mini Projects	170
A.0 About This Appendix	171
A.1 The Running Scenario: CityPulse	171
A.2 Laboratory Setup.....	172
Section A – Laboratory Experiments	173
Section B – Tool-Wise Assignments	174
Section C – Industry Case Studies	176
Section D – Mini Projects	178
Section E – Report Format, Rubrics and Viva Bank.....	179