

TABLE OF CONTENTS

Part I: Foundations

Chapter 1: Why NoSQL? The Four Families and the RDBMS Contrast.....	1
Learning Objectives.....	1
Concept Introduction: The Rise of NoSQL.....	1
The Four NoSQL Families.....	2
Key-Value Stores.....	2
Document Databases.....	3
Column-Family (Wide-Column) Stores.....	3
Graph Databases.....	4
RDBMS vs NoSQL: A Structured Comparison.....	4
ACID vs BASE: Two Philosophies of Consistency.....	5
ACID Properties.....	5
BASE Properties.....	6
The CAP Theorem.....	6
Survey of NoSQL and Cloud-Native Database Products.....	7
Best Practices and Production Examples.....	8
Hands-On Labs.....	9
[Beginner] Lab 1.1: Classify Database Products.....	9
Solution: Lab 1.1.....	9
[Intermediate] Lab 1.2: Map a Real-World Application to NoSQL Families.....	10
Solution: Lab 1.2.....	10
Chapter Summary.....	11
Chapter 2: Cloud Computing Primer for Database Engineers	12
Learning Objectives.....	12
Concept Introduction: What Is Cloud Computing?.....	12
IaaS: Infrastructure as a Service.....	12
PaaS: Platform as a Service.....	13
SaaS: Software as a Service.....	13
DBaaS: Database as a Service.....	13
Comparing Major Cloud Providers.....	14

Installation and Environment Setup.....	14
Setting Up MongoDB Atlas (Cloud).....	15
Local MongoDB via Docker.....	15
Installing PyMongo.....	16
Code Walkthrough: Your First MongoDB Connection.....	16
Installing and Using MongoDB Compass.....	19
Atlas CLI Basics.....	19
Best Practices and Production Examples.....	20
Hands-On Labs.....	20
[Beginner] Lab 2.1: Set Up Your Development Environment.....	20
Solution: Lab 2.1.....	21
[Beginner] Lab 2.2: Explore Database Operations with PyMongo.....	21
Solution: Lab 2.2.....	22
Chapter Summary.....	24
Chapter 3: MongoDB Data Model Fundamentals.....	25
Learning Objectives.....	25
Concept Introduction: Documents, Collections, and BSON.....	25
BSON: Binary JSON.....	25
Documents and Collections.....	27
Embedding vs Referencing.....	27
Embedding: The One-to-Few Pattern.....	28
Referencing: The One-to-Many and Many-to-Many Patterns.....	29
Design Methods: When to Embed vs Reference.....	31
Schema Validation.....	32
Best Practices and Production Examples.....	34
Hands-On Labs.....	35
[Intermediate] Lab 3.1: Design an Embedded Document Schema.....	35
Solution: Lab 3.1.....	35
[Intermediate] Lab 3.2: Migrate from Referenced to Embedded Design.....	35
Solution: Lab 3.2.....	36
Chapter Summary.....	36
Chapter 4: Mastering Queries: CRUD and the Query Language in Python.....	37
Learning Objectives.....	37

Concept Introduction: The MongoDB Query Language.....	37
Create Operations.....	38
Read Operations and Query Operators.....	40
Comparison Operators.....	40
Logical Operators.....	41
Array Operators.....	42
Projections and Cursor Operations.....	44
Update Operations.....	45
Delete Operations.....	47
Best Practices and Production Examples.....	47
Hands-On Labs.....	47
[Intermediate] Lab 4.1: Build a Product Query API.....	47
Solution: Lab 4.1.....	48
[Intermediate] Lab 4.2: Implement Upsert and Bulk Operations.....	48
Solution: Lab 4.2.....	49
Chapter Summary.....	49
Chapter 5: Aggregation Framework & Advanced Querying.....	50
Learning Objectives.....	50
Concept Introduction: Why Aggregation?.....	50
Core Aggregation Stages.....	51
\$match: Filtering in a Pipeline.....	51
\$group: Aggregating Data.....	51
\$project: Reshaping Documents.....	53
\$unwind: Deconstructing Arrays.....	53
\$facet: Multi-Faceted Aggregation.....	54
Pipeline Optimization Strategies.....	55
Best Practices and Production Examples.....	55
Hands-On Labs.....	56
[Intermediate] Lab 5.1: Build a Sales Analytics Pipeline.....	56
Solution: Lab 5.1.....	56
[Intermediate] Lab 5.2: Join Data with \$lookup.....	56
Solution: Lab 5.2.....	56
Chapter Summary.....	57
Chapter 6: Indexing Strategies & Performance Tuning	58

Learning Objectives.....	58
Concept Introduction: Why Indexes Matter.....	58
Index Types.....	59
Single-Field Indexes.....	59
Compound Indexes and the ESR Rule.....	61
Multikey Indexes.....	62
Text Indexes and Atlas Search.....	62
Using explain () to Diagnose Performance.....	64
Best Practices and Production Examples.....	65
Hands-On Labs.....	65
[Intermediate] Lab 6.1: Index Design for a Query Workload.....	65
Solution: Lab 6.1.....	65
[Advanced] Lab 6.2: Compare Index Strategies.....	66
Solution: Lab 6.2.....	67
Chapter Summary.....	67

Part II: Architecture, Security & Operations

Chapter 7: MongoDB Architecture Deep Dive.....	69
Learning Objectives.....	69
Concept Introduction: MongoDB Internals.....	69
The Wired Tiger Storage Engine.....	69
Canonical MongoDB Use Cases.....	73
Normalized vs. Embedded Design Revisited.....	73
Best Practices and Production Examples.....	73
Hands-On Labs.....	74
[Intermediate] Lab 7.1: Analyze a Production Schema.....	74
Solution: Lab 7.1.....	74
[Advanced] Lab 7.2: Working Set Analysis.....	75
Solution: Lab 7.2.....	75
Chapter Summary.....	75
Chapter 8: Replication & High Availability.....	76
Learning Objectives.....	76
Concept Introduction: Why Replication?.....	76

Replica Set Configuration.....	77
Elections and Failover.....	78
Read and Write Concerns.....	80
Best Practices and Production Examples.....	81
Hands-On Labs.....	82
[Advanced] Lab 8.1: Deploy a Local Replica Set.....	82
Solution: Lab 8.1.....	82
[Advanced] Lab 8.2: Observe Failover and Election.....	82
Solution: Lab 8.2.....	83
Chapter Summary.....	83
Chapter 9: CAP Theorem in Practice.....	84
Learning Objectives.....	84
Concept Introduction: CAP Beyond the Theorem.....	84
Write Concern Trade-Offs.....	85
Read Concern Trade-Offs.....	86
Worked Scenarios.....	87
Scenario 1: Social Media Feed.....	87
Scenario 2: Banking System.....	87
Best Practices and Production Examples.....	87
Hands-On Labs.....	88
[Advanced] Lab 9.1: Measure Write Concern Latency.....	88
Solution: Lab 9.1.....	88
[Advanced] Lab 9.2: Simulate a Network Partition.....	89
Solution: Lab 9.2.....	89
Chapter Summary.....	90
Chapter 10: Transactions, Atomicity, Durability & Journaling.....	91
Learning Objectives.....	91
Concept Introduction: Atomicity at Two Levels.....	91
Multi-Document Transactions.....	93
Journaling and Crash Recovery.....	94
Best Practices and Production Examples.....	94
Hands-On Labs.....	95
[Advanced] Lab 10.1: Implement a Bank Transfer Transaction.....	95

Solution: Lab 10.1.....	95
[Advanced] Lab 10.2: Inventory Reservation with Deadlock Handling.....	96
Solution: Lab 10.2.....	96
Chapter Summary.....	97
Chapter 11: Sharding & Horizontal Scalability.....	98
Learning Objectives.....	98
Concept Introduction: Scaling Beyond a Single Server.....	98
Shard Key Selection.....	99
The Balancer and Chunk Migration.....	101
Best Practices and Production Examples.....	101
Hands-On Labs.....	101
[Intermediate] Lab 11.1: Evaluate Shard Key Choices.....	101
Solution: Lab 11.1.....	102
[Intermediate] Lab 11.2: Shard Key Hotspot Detection	102
Solution: Lab 11.2.....	102
Chapter Summary.....	103
Chapter 12: Database Security & the CIA Triad.....	104
Learning Objectives.....	104
Concept Introduction: The CIA Triad for Databases.....	104
Confidentiality.....	104
Integrity	105
Availability.....	105
Threat Modeling Exercise.....	106
Best Practices and Production Examples.....	108
Hands-On Labs.....	108
[Advanced] Lab 12.1: Configure Authentication and RBAC.....	108
Solution: Lab 12.1.....	108
[Intermediate] Lab 12.2: TLS/Encryption-in-Transit Setup	109
Solution: Lab 12.2.....	109
Chapter Summary.....	110
Part III: Production & AI-Native Platforms	
Chapter 13: Deploying MongoDB with Containers and the Cloud.....	111
Learning Objectives.....	111

Concept Introduction: From Local to Production.....	111
Docker and Docker Compose for MongoDB.....	111
Kubernetes Basics for MongoDB.....	113
Infrastructure as Code with Terraform.....	115
Best Practices and Production Examples.....	116
Hands-On Labs.....	116
[Advanced] Lab 13.1: Deploy a Replica Set with Docker Compose.....	116
Solution: Lab 13.1.....	116
[Advanced] Lab 13.2: Kubernetes StatefulSet for MongoDB.....	117
Solution: Lab 13.2.....	117
Chapter Summary.....	118
Chapter 14: Managing Database Services in Production.....	119
Learning Objectives.....	119
Concept Introduction: The Operational Side of Databases.....	119
Monitoring and Alerting.....	121
Backup Strategies and Restore Drills.....	122
Disaster Recovery: RPO and RTO.....	122
Capacity Planning.....	122
Best Practices and Production Examples.....	122
Hands-On Labs.....	123
[Advanced] Lab 14.1: Build a Monitoring Dashboard.....	123
Solution: Lab 14.1.....	123
[Intermediate] Lab 14.2: Automated Backup and Restore Verification	124
Solution: Lab 14.2.....	124
Chapter Summary.....	125
Chapter 15: Designing and Architecting a MongoDB-Backed Product as a Service.....	126
Learning Objectives.....	126
Concept Introduction: Database as a Service.....	126
Multi-Tenant Schema Design.....	126
API Layer and Rate Limiting.....	130
Polyglot Persistence.....	131

Best Practices and Production Examples.....	131
Hands-On Labs.....	132
[Advanced] Lab 15.1: Design a Multi-Tenant API.....	132
Solution: Lab 15.1.....	132
[Advanced] Lab 15.2: Data Migration Between Tenant Isolation Strategies	132
Solution: Lab 15.2.....	133
Chapter Summary.....	134
Chapter 16: Wide-Column and Graph Databases: Cassandra and Neo4j.....	135
Learning Objectives.....	135
Concept Introduction: Beyond Document Databases.....	135
Apache Cassandra: Column-Family Model.....	135
Neo4j: Graph Database Fundamentals.....	136
When to Use Each Database.....	137
Hands-On Labs.....	140
[Intermediate] Lab 16.1: Model Data in Three Databases.....	140
Solution: Lab 16.1.....	140
[Intermediate] Lab 16.2: Build a Hybrid Query Service	140
Solution: Lab 16.2.....	141
Chapter Summary.....	142
Chapter 17: AI-Native MongoDB: Vector Search and RAG Foundations.....	144
Learning Objectives.....	144
Concept Introduction: MongoDB as an AI-Native Data Platform.....	144
Embeddings and Vector Representations.....	145
Atlas Vector Search.....	145
Building a RAG Pipeline.....	146
GraphRAG: Beyond Vector Search.....	148
Best Practices and Production Examples.....	149
Hands-On Labs.....	150
[Advanced] Lab 17.1: Build a Simple RAG Pipeline.....	150
Solution: Lab 17.1.....	150
[Advanced] Lab 17.2: Compare HNSW and IVF Vector Index Types ...	151
Solution: Lab 17.2.....	151
Chapter Summary.....	152

Chapter 18: Comparative Capstone Lab	153
Learning Objectives.....	153
Concept Introduction: Putting It All Together.....	153
The Use Case: Financial Fraud Detection.....	153
Schema Design Comparison.....	154
MongoDB Schema.....	154
Cassandra Schema.....	154
Neo4j Schema.....	154
Query Comparison.....	154
Comparative Write-Up Template.....	156
Best Practices and Production Examples.....	156
Hands-On Labs.....	158
[Advanced] Lab 18.1: Build the Comparative System.....	158
Solution: Lab 18.1.....	158
[Advanced] Lab 18.2: Write a Comparative Analysis Report	159
Solution: Lab 18.2.....	159
Chapter Summary.....	160
 Appendix A: Building and Deploying an LLM-Backed Chat Application on MongoDB	161
 Appendix B: Answer Key for Self-Check Questions	168
 Glossary of Terms	178